
BY COMMUNITY • FOR COMMUNITY •
#GLOBALAZURE
TORINO



Azure Meetup
TORINO

alveo®
EXPERTISE

TD SYNNEX

#GlobalAzureTorino

Mastering AI Multi-Agent Cyber Warfare: Dall'Attacco alla Difesa Adattiva

Nino Crudele



Azure Meetup TORINO

alveo EXPERTISE

TD SYNNEX

reverie.herokuapp.com/arXiv_Demo/
X3M.AI | GlobalAzure | Parasol TRIA | MITRE | TO WATCH&READ | SGN | Jobs | Install NVIDIA GPU... | linkedin sth | MI5 | meetup | SIS | Science and Te... | Jobs - Recruitment...



This is a pre-computed replay of a simulation that accompanies the paper entitled "Generative Agents: Interactive Simulacra of Human Behavior." It is for demonstration purposes only.

Current Time:

Gli agents si sono evoluti in modi inaspettati

Abstract:

The promise of human behavioral simulation—general-purpose computational agents that replicate human behavior across domains—could enable broad applications in policymaking and social science. We present a novel agent architecture that simulates the attitudes and behaviors of 1,052 real individuals—applying large language models to qualitative interviews about their lives, then measuring how well these agents replicate the attitudes and behaviors of the individuals that they represent. The generative agents replicate participants' responses on the General Social Survey 85% as accurately as participants replicate their own answers two weeks later, and perform comparably in predicting personality traits and outcomes in experimental replications. Our architecture reduces accuracy biases across racial and ideological groups compared to agents given demographic descriptions. This work provides a foundation for new tools that can help investigate individual and collective behavior.





Human Participants

Simulations



2-hr Audio Interview
(Avg. 6,491 words)

Interview script drawn from the American Voices Project

Generative Agents

Interview transcript serves as agent memory

Actual participant responses

General Social Survey (177 Items)
Big Five Personality Inventory (44 Items)
Economic Games (5 Items)
Behavioral Experiments (5 Items)

Simulated participant responses

General Social Survey (177 Items)
Big Five Personality Inventory (44 Items)
Economic Games (5 Items)
Behavioral Experiments (5 Items)

Compare actual to simulated responses, adjusting for participant self-consistency

#GlobalAzureTorino

Chi sono



- Hacker (a.k.a 90)
- Founder of X3M.AI
  - Cybersecurity Critica per Governi, Associazioni e Aziende Strategiche.
  - Assessment e Audit di infrastrutture ad alta sensibilità.
  - Cyber Intelligence offensiva e difensiva su minacce avanzate.
  - Supporto operativo per enti governativi, grandi aziende e organizzazioni strategiche.
- 18x Microsoft MVP, Certified Ethical Hacher, Master in Cybersecurity

- https://www.linkedin.com/in/ninocrudele/

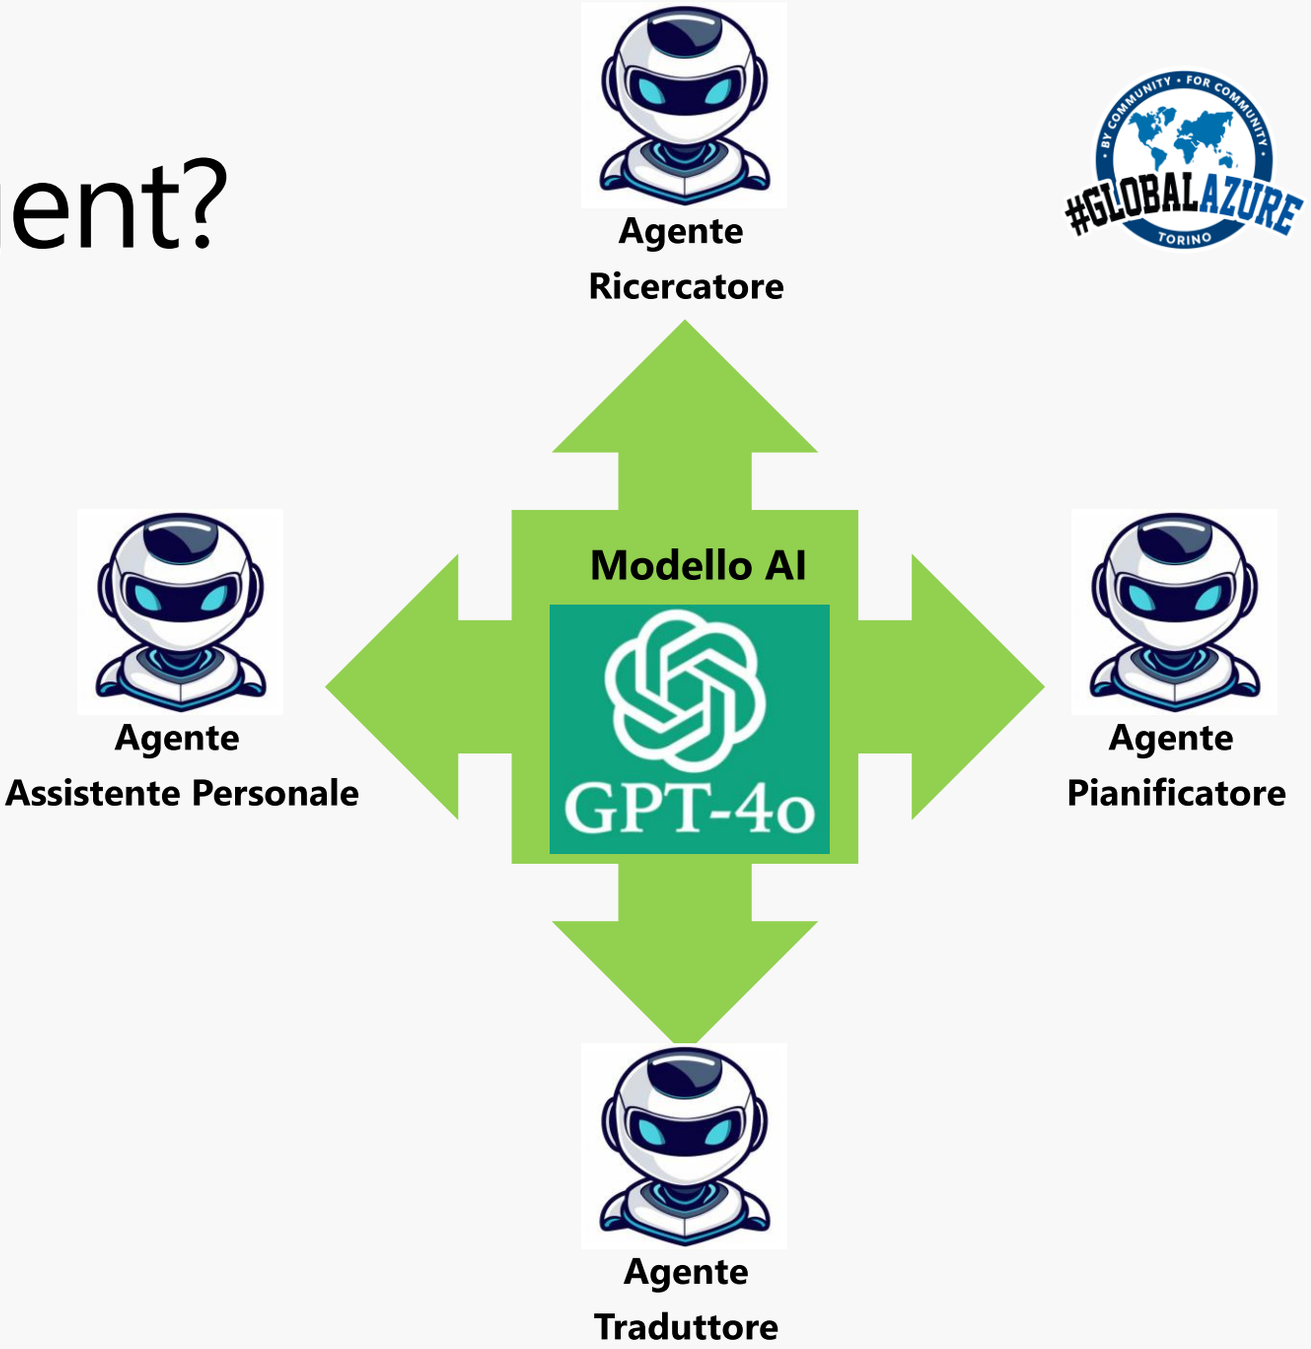
MVP Microsoft® Most Valuable Professional

C|EH CERTIFIED ETHICAL HACKER

#GlobalAzureTorino

Che cosa e' un AI Agent?

Un agente AI è un programma intelligente specializzato che può osservare cosa succede, capire la situazione e decidere cosa fare da solo, per raggiungere un obiettivo.



La Cyber Warfare si e’ evoluta







cybermap.kaspersky.com

CYBERTHREAT LIVE MAP

EN

Protect Yourself

MAP STATISTICS DATA SOURCES BUZZ WIDGET

Share

BHUTAN

157 MOST-ATTACKED COUNTRY

OAS	8
ODS	89
MAV	0
WAV	12
IDS	0
VUL	0
KAS	520
BAD	0
[illegible]	0

Detections discovered since 00:00 GMT

More details

Share data



MAV - Mail Anti Virus

MAV (Mail Anti-Virus) shows malware detection flow during Mail Anti-Virus scan when new objects appear in an email application (Outlook, The Bat, Thunderbird). The MAV scans incoming messages and calls OAS when saving attachments to a disk.

2979680	2102809	372527	1760747	823340	16571	5191039	0	67024
OAS	ODS	MAV	WAV	IDS	VUL	KAS	BAD	[illegible]

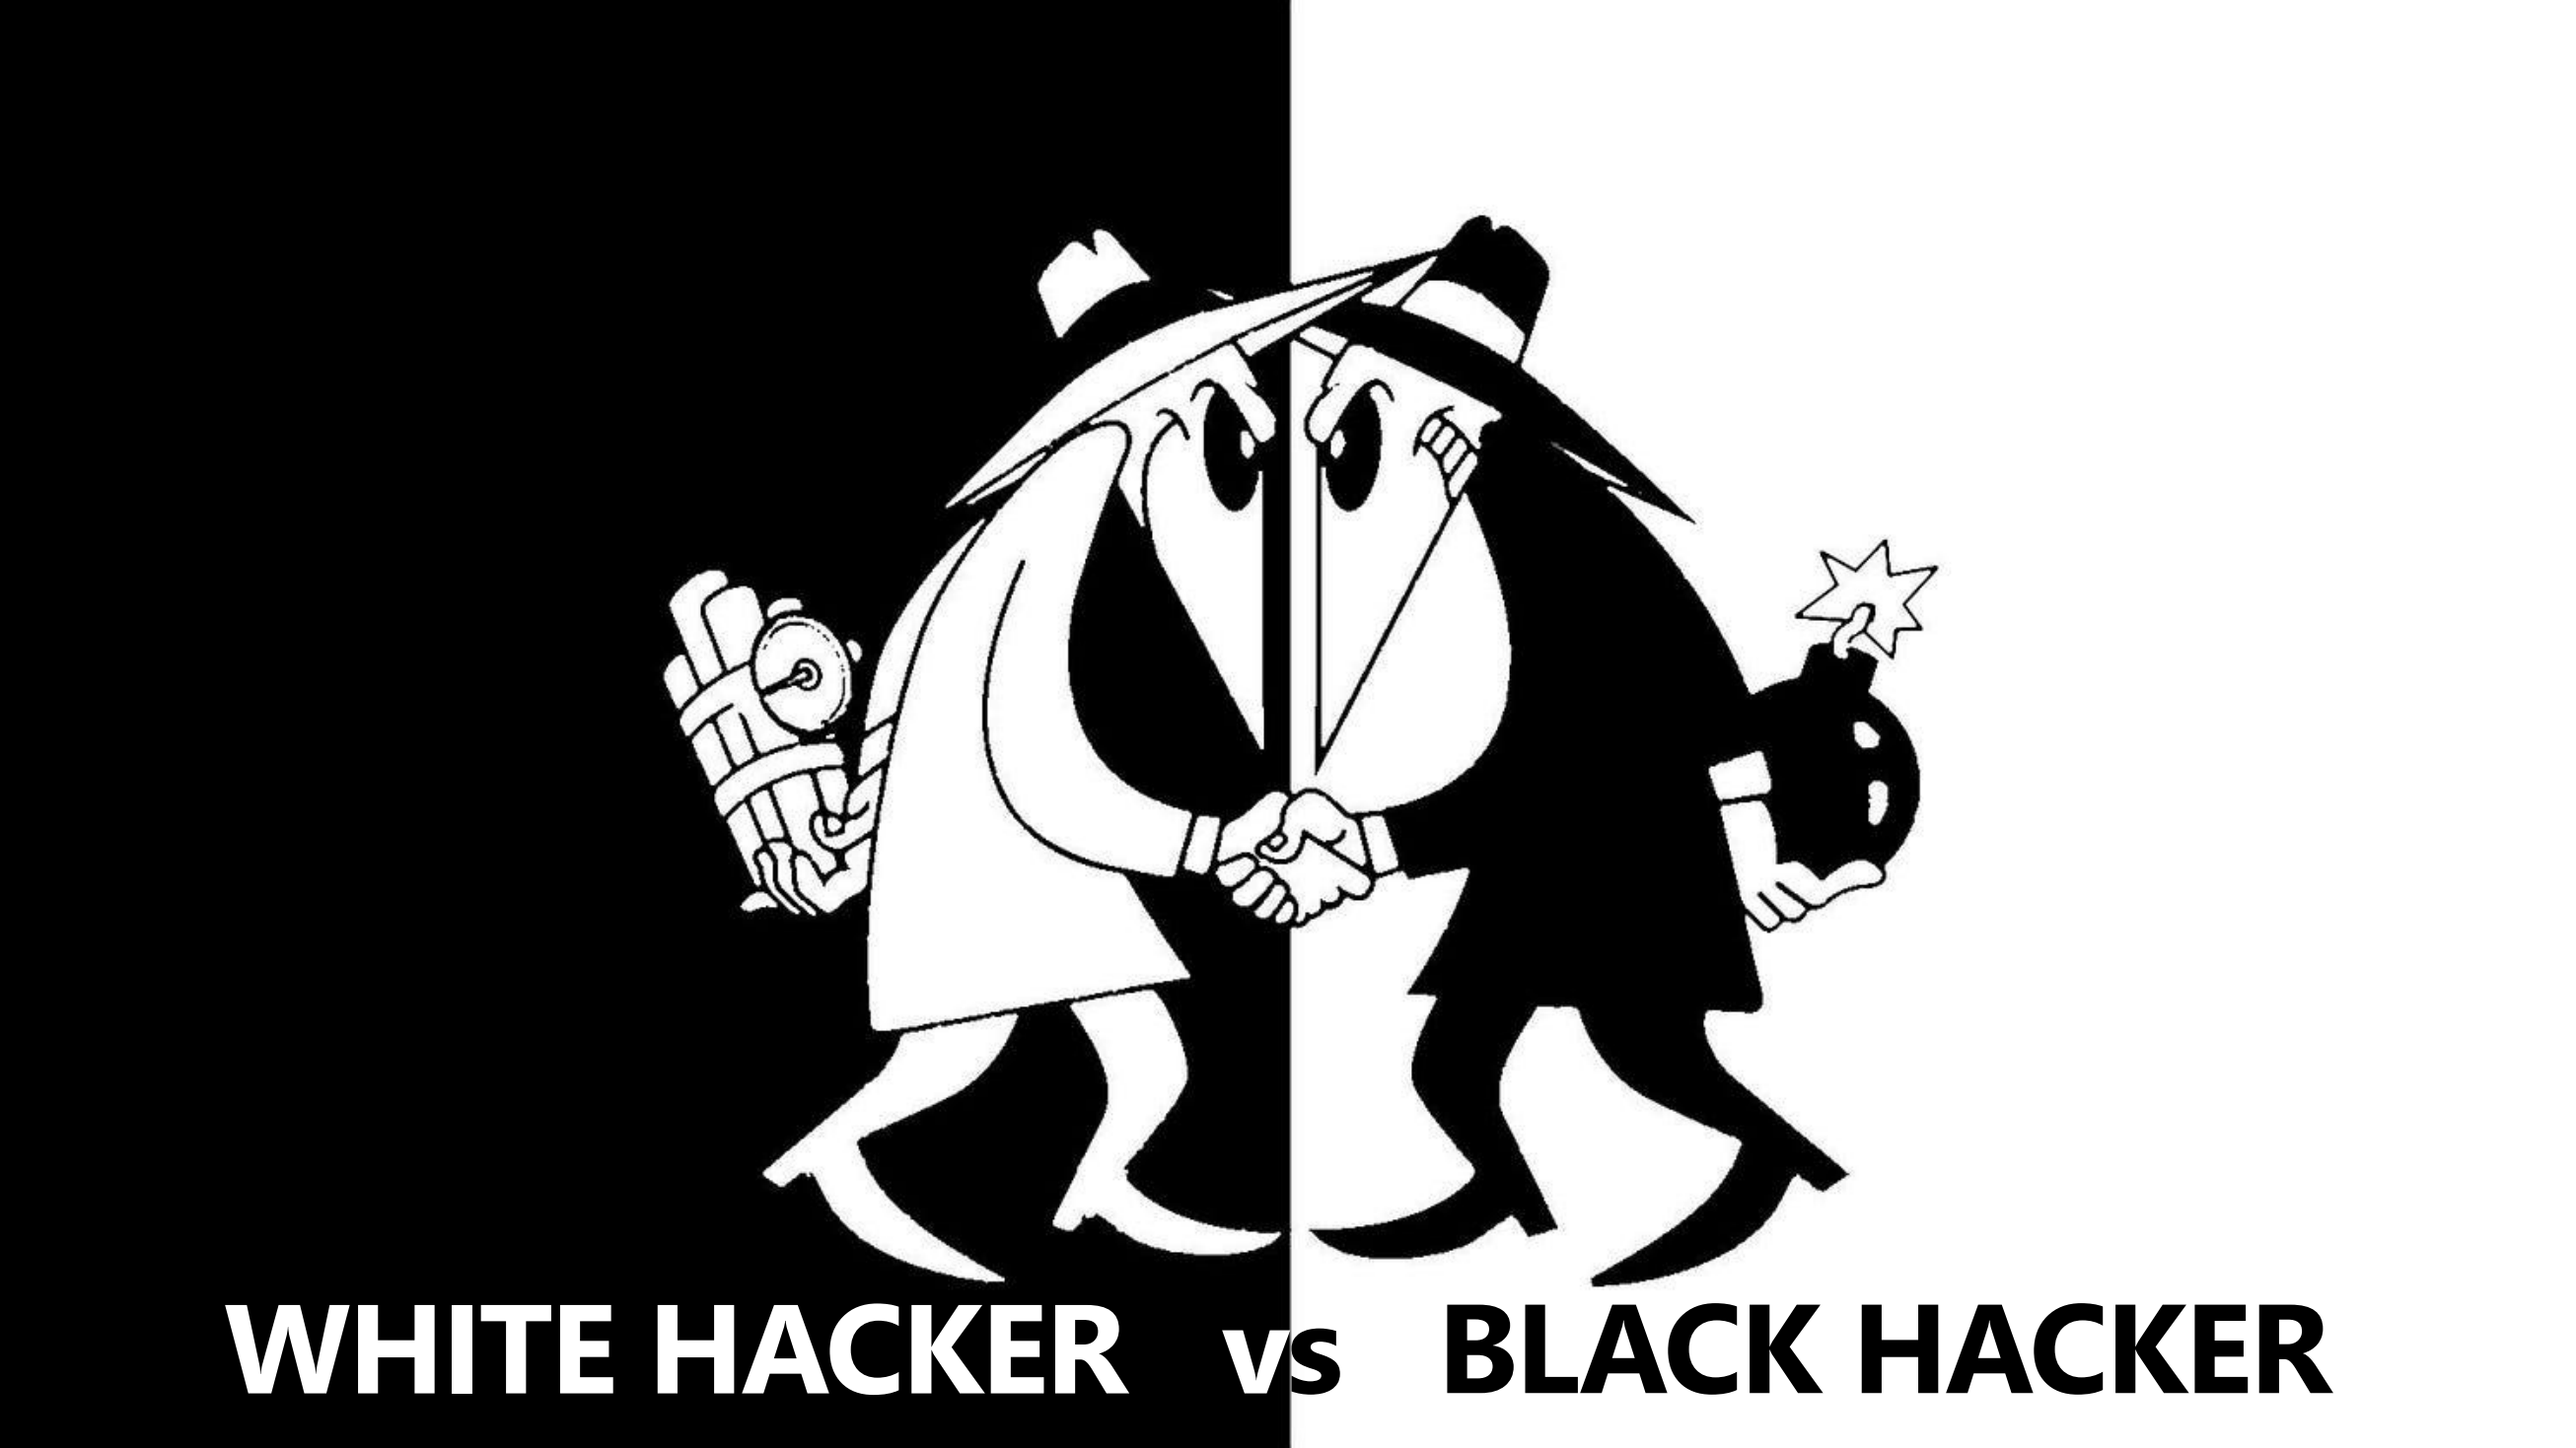
kaspersky

Based on data from Kaspersky.

© 2025 AO Kaspersky Lab. All Rights Reserved.
Terms of Service Privacy policy Cookies

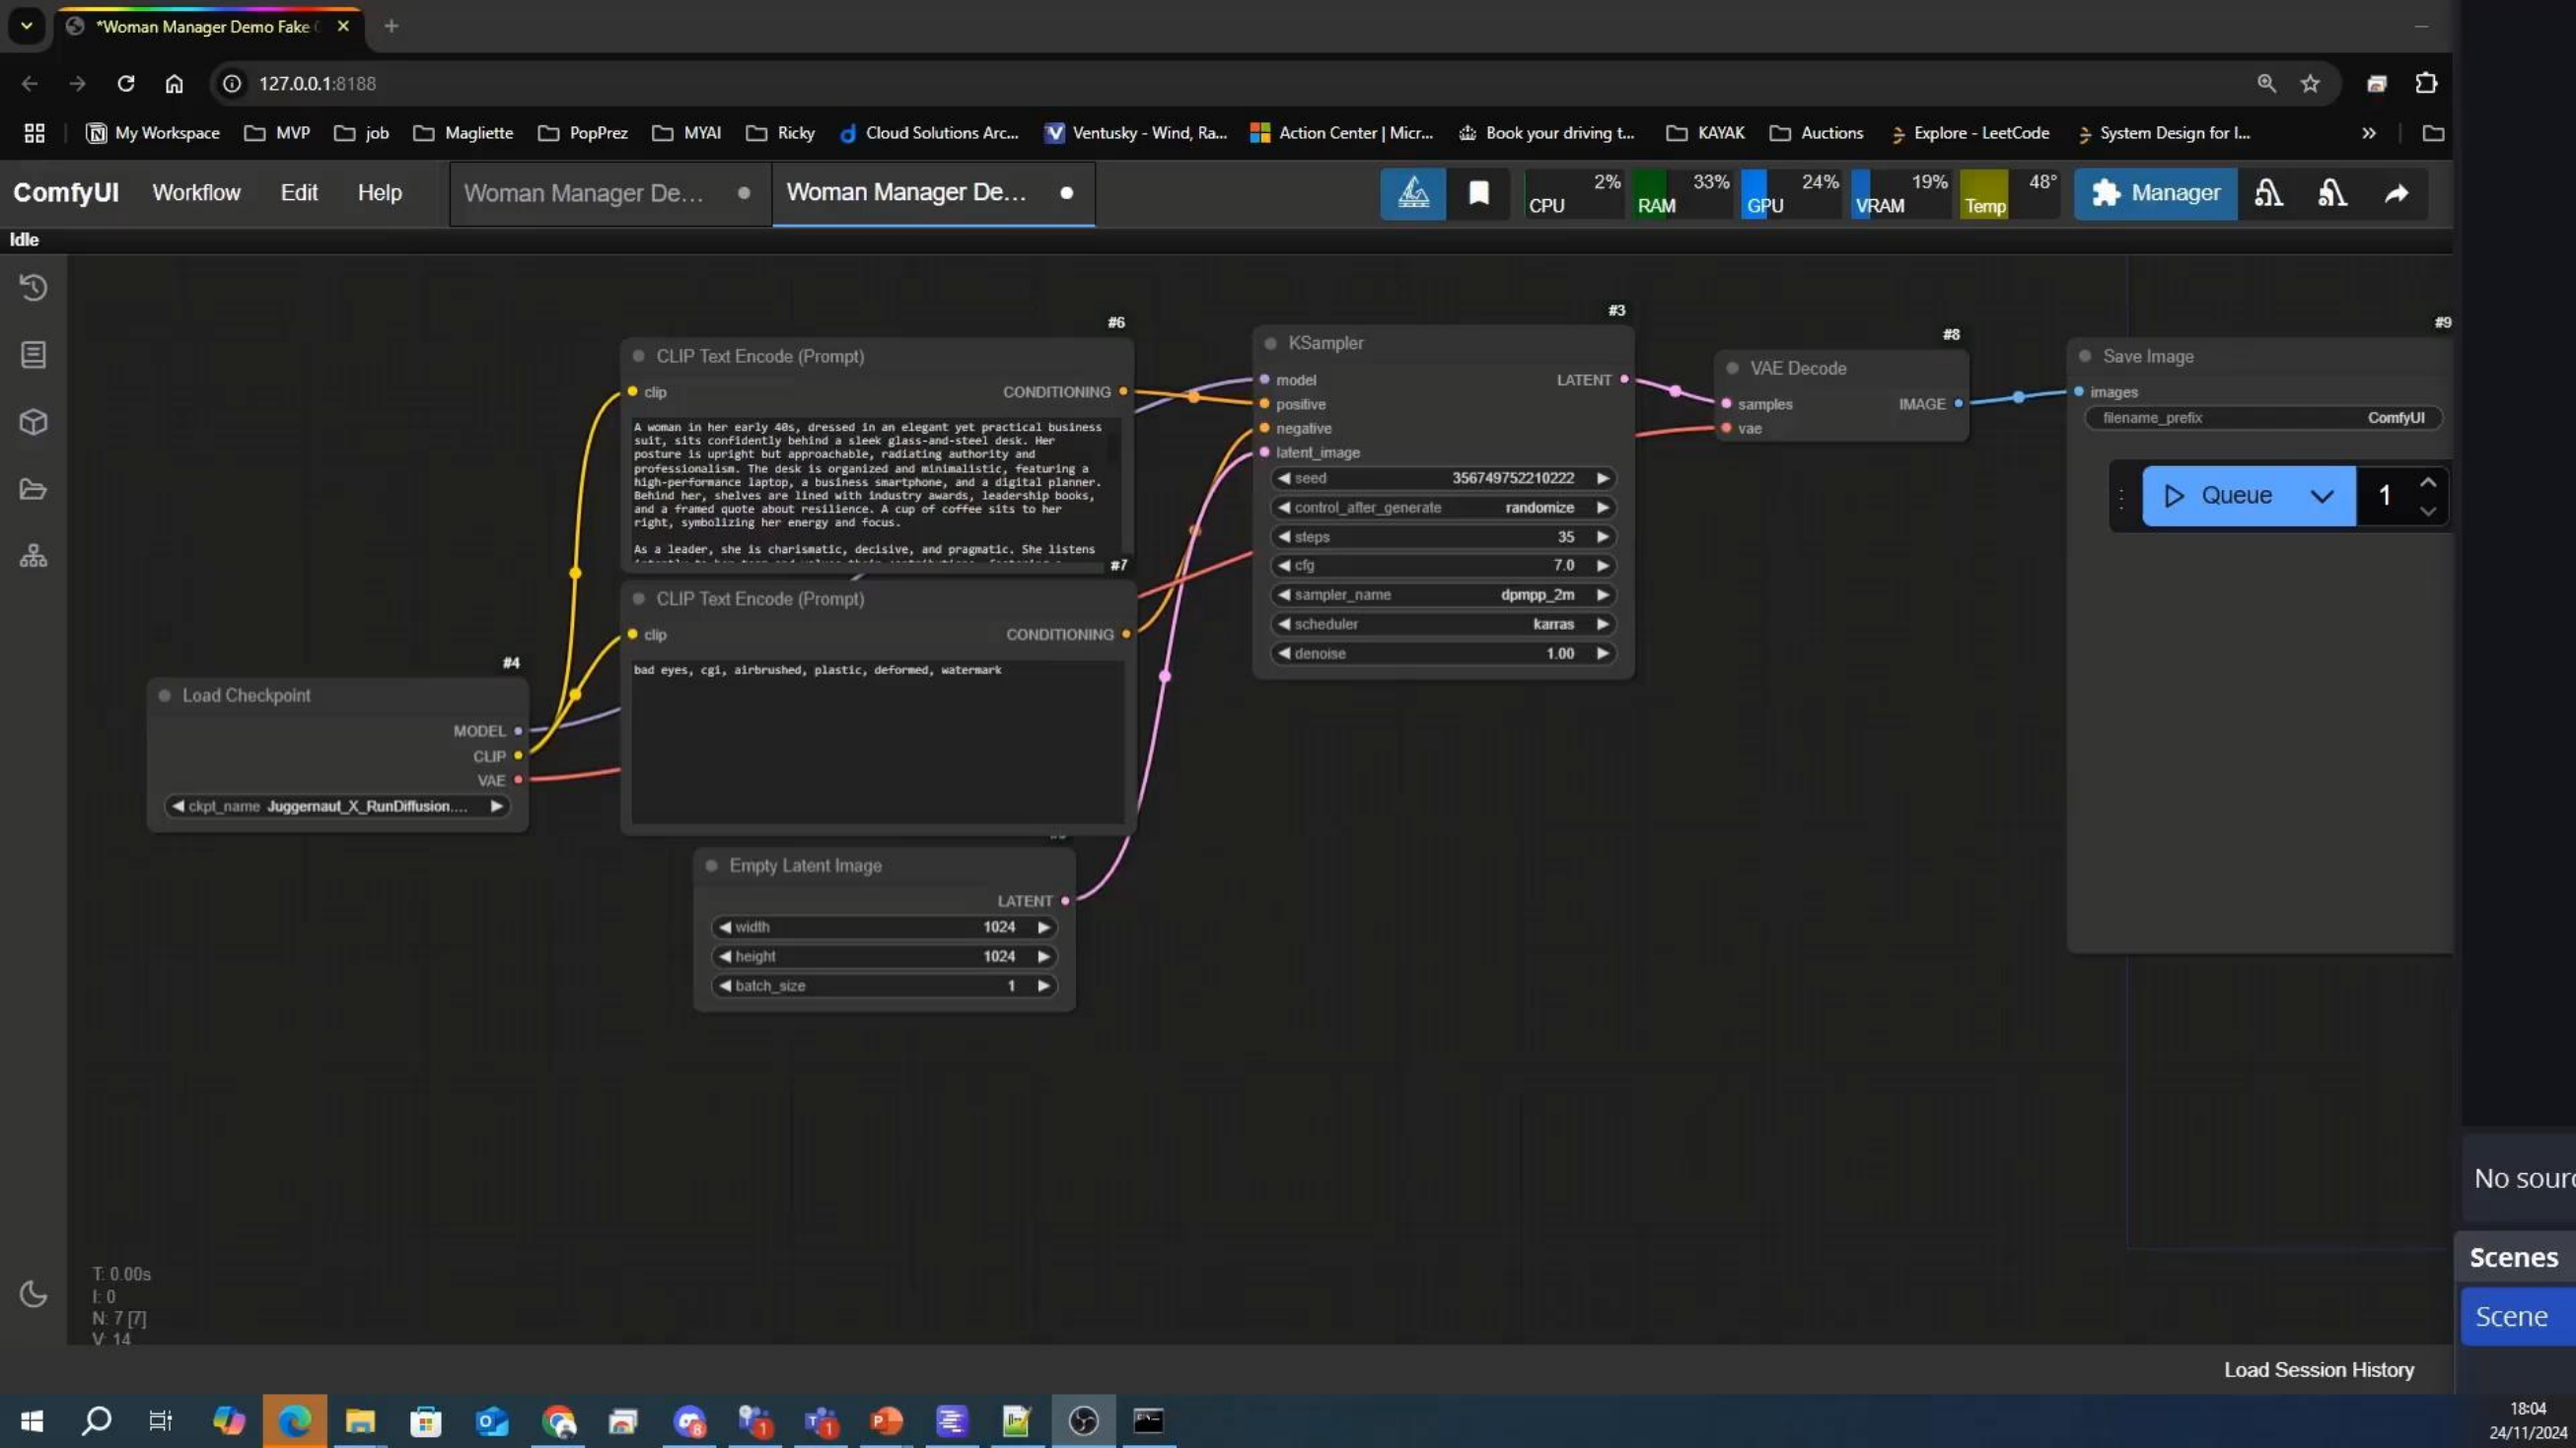
Search

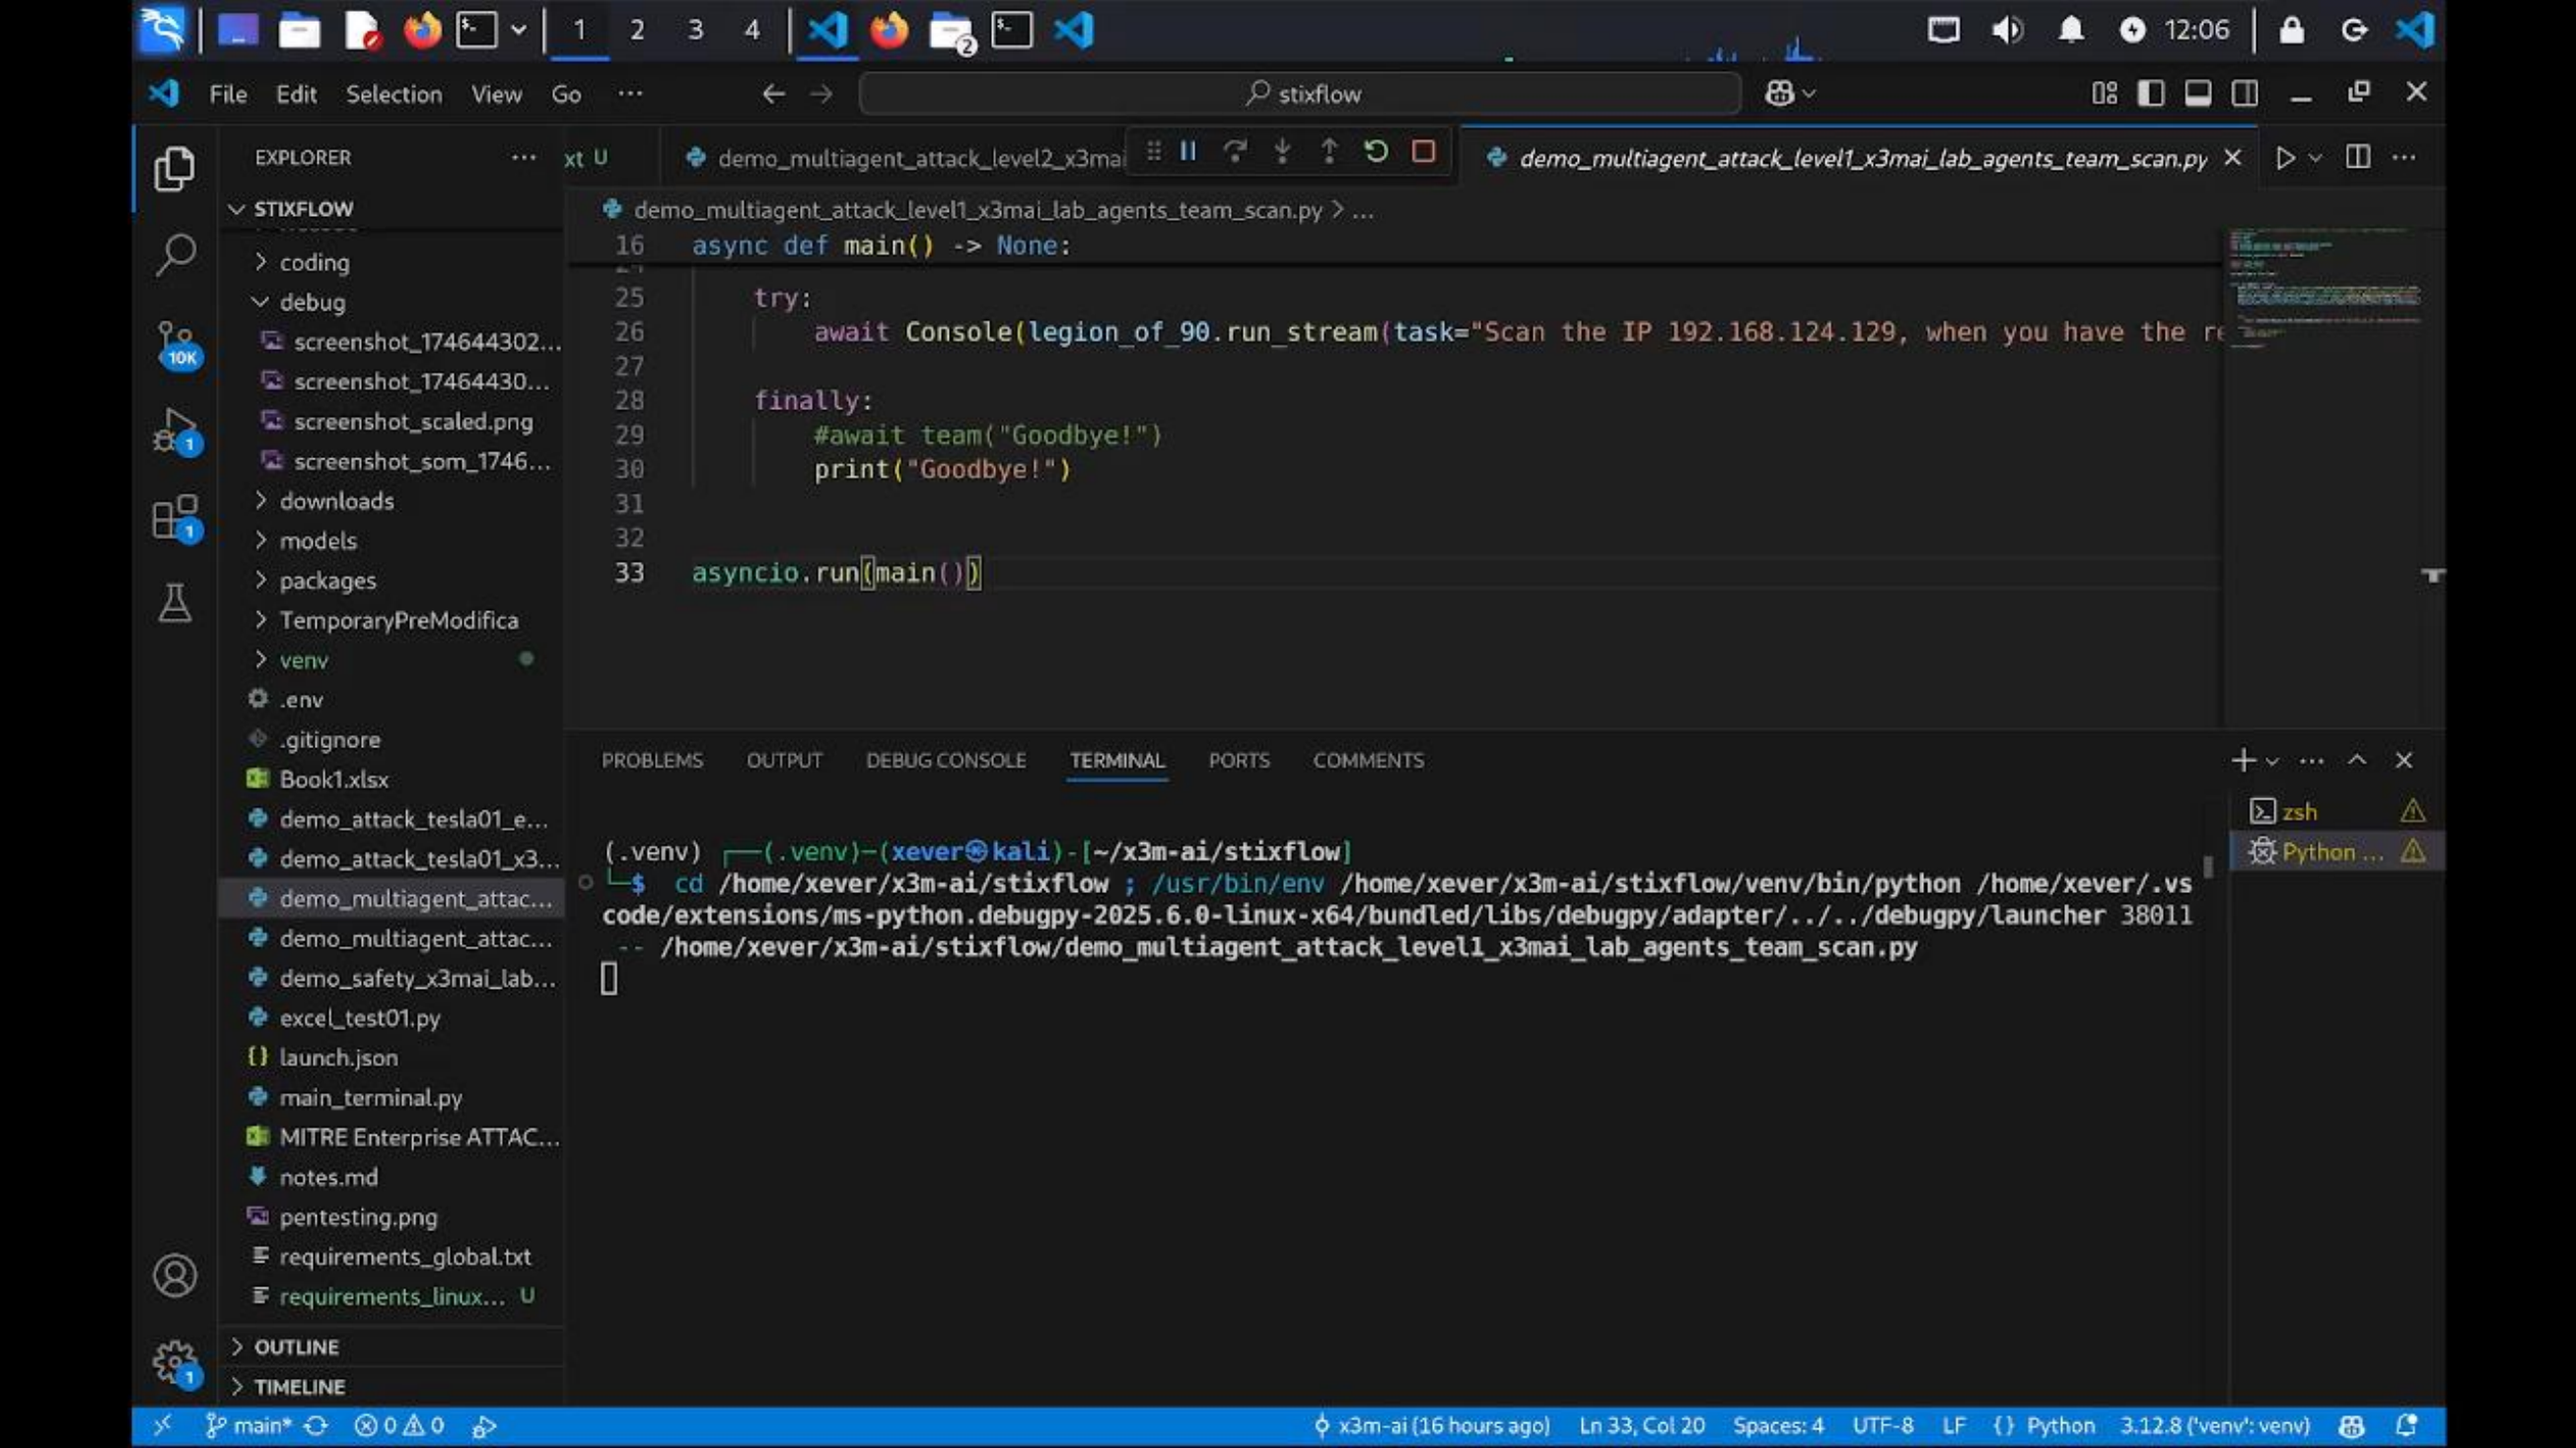
12:56
05/05/2025


WHITE HACKER vs BLACK HACKER



Il lato oscuro di AI


*Woman Manager Demo Fake
127.0.0.1:8188
My Workspace
MVP
job
Magliette
PopPrez
MYAI
Ricky
Cloud Solutions Arc...
Ventusky - Wind, Ra...
Action Center | Micr...
Book your driving t...
KAYAK
Auctions
Explore - LeetCode
System Design for I...
ComfyUI
Workflow
Edit
Help
Woman Manager De...
Woman Manager De...
CPU 2%
RAM 33%
GPU 24%
VRAM 19%
Temp 48°
Manager
Idle
#6
CLIP Text Encode (Prompt)
clip
CONDITIONING
A woman in her early 40s, dressed in an elegant yet practical business suit, sits confidently behind a sleek glass-and-steel desk. Her posture is upright but approachable, radiating authority and professionalism. The desk is organized and minimalistic, featuring a high-performance laptop, a business smartphone, and a digital planner. Behind her, shelves are lined with industry awards, leadership books, and a framed quote about resilience. A cup of coffee sits to her right, symbolizing her energy and focus.
As a leader, she is charismatic, decisive, and pragmatic. She listens
#7
CLIP Text Encode (Prompt)
clip
CONDITIONING
bad eyes, cgi, airbrushed, plastic, deformed, watermark
#4
Load Checkpoint
MODEL
CLIP
VAE
ckpt_name Juggernaut_X_RunDiffusion....
Empty Latent Image
LATENT
width 1024
height 1024
batch_size 1
#3
KSampler
model
positive
negative
latent_image
LATENT
seed 356749752210222
control_after_generate randomize
steps 35
cfg 7.0
sampler_name dpmpp_2m
scheduler karras
denoise 1.00
#8
VAE Decode
samples
vae
IMAGE
#9
Save Image
images
filename_prefix ComfyUI
Queue
1
T: 0.00s
I: 0
N: 7 [7]
V: 14
Load Session History
No sour
Scenes
Scene
18:04
24/11/2024


File Edit Selection View Go
stixflow
EXPLORER
STIXFLOW
coding
debug
screenshot_174644302...
screenshot_17464430...
screenshot_scaled.png
screenshot_som_1746...
downloads
models
packages
TemporaryPreModifica
venv
.env
.gitignore
Book1.xlsx
demo_attack_tesla01_e...
demo_attack_tesla01_x3...
demo_multiagent_attac...
demo_multiagent_attac...
demo_safety_x3mai_lab...
excel_test01.py
launch.json
main_terminal.py
MITRE Enterprise ATTAC...
notes.md
pentesting.png
requirements_global.txt
requirements_linux... U
OUTLINE
TIMELINE
demo_multiagent_attack_level2_x3mai
demo_multiagent_attack_level1_x3mai_lab_agents_team_scan.py
demo_multiagent_attack_level1_x3mai_lab_agents_team_scan.py > ...
16 async def main() -> None:
25 try:
26 await Console(legion_of_90.run_stream(task="Scan the IP 192.168.124.129, when you have the re
27
28 finally:
29 #await team("Goodbye!")
30 print("Goodbye!")
31
32
33 asyncio.run(main())
PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS COMMENTS
zsh
Python ...
(.venv) ┌──(.venv)─(xever㉿kali)-[~/x3m-ai/stixflow]
└─$ cd /home/xever/x3m-ai/stixflow ; /usr/bin/env /home/xever/x3m-ai/stixflow/venv/bin/python /home/xever/.vscode/extensions/ms-python.debugpy-2025.6.0-linux-x64/bundled/libs/debugpy/adapter/../../debugpy/launcher 38011 -- /home/xever/x3m-ai/stixflow/demo_multiagent_attack_level1_x3mai_lab_agents_team_scan.py
main* 0 0
x3m-ai (16 hours ago)
Ln 33, Col 20
Spaces: 4
UTF-8
LF
Python
3.12.8 ('venv': venv)

I Guardiani della Nuova Intelligenza



Create a resource
Home
Dashboard
All services
FAVORITES
All resources
Resource groups
App Services
Function App
SQL databases
Azure Cosmos DB
Virtual machines
Load balancers
Storage accounts
Virtual networks
Microsoft Entra ID
Monitor
Advisor
Microsoft Defender for Cloud
Cost Management + Billing
Help + support

Purview IRM ('58b3ee83') Data Exfiltration Monitoring involving one user

Incident number 4251

Refresh | Logs | Tasks | Activity log

This is the new, improved incident page - **Now generally available**. You can use the toggle to switch back. New experience

Write permissions on the workspace are required to modify incidents

High	New	Unassigned
Severity	Status	Owner

Investigate in Microsoft Defender XDR

Incident summary

AI-generated content may be incorrect. Check it for accuracy

Last Updated: 5/5/2025, 3:15:34 PM

The high severity incident "Purview IRM ('58b3ee83') Data Exfiltration Monitoring involving one user" occurred between 2025-03-07 14:34:35 UTC and 2025-05-05 12:21:38 UTC. The incident impacted the user 'jonaw'.

- The incident began at 2025-03-07 14:34:35 UTC with a high severity alert for data exfiltration monitoring by Purview IRM ('58b3ee83'). This alert indicated potential data exfiltration activities involving the account **'jonaw'** (jonaw@vnevado.alpineskihouse.co) (**Exfiltration** killchain stage).

Workspace name
cybersoc

The investigation graph requires that your incident includes entities (for example: user, host, IP, etc.). Use the entity mapping option when defining your alerts. Learn more >

Overview | Entities

Incident actions

Incident timeline

Search | Add filter

Mar 7 14:34:35	Purview IRM ('58b3ee83') Data Exf... Low | Detected by Microsoft 36... | Ta...
Mar 7 14:34:35	Purview IRM ('58b3ee83') Data Exf... Low | Detected by Microsoft 36... | Ta...
Mar 7 14:34:35	Purview IRM ('58b3ee83') Data Exf... Hi... | Detected by Microsoft 36... | Ta...
Mar 7 14:34:35	Purview IRM ('58b3ee83') Data Exf...

Entities

No Entities

No entities found

Similar incidents

No similar incidents to display

Come opera sul campo



Il futuro con AI – cosa fare?

- Impara il Prompt Engineering
  - Perche' e' il modo di lavorare con AI
- Impara Python
  - Perche' e' il linguaggio utilizzato da AI
- Esperimenta il piu' possible!
  - Perche' e' il business del futuro

Ma prima di tutto…

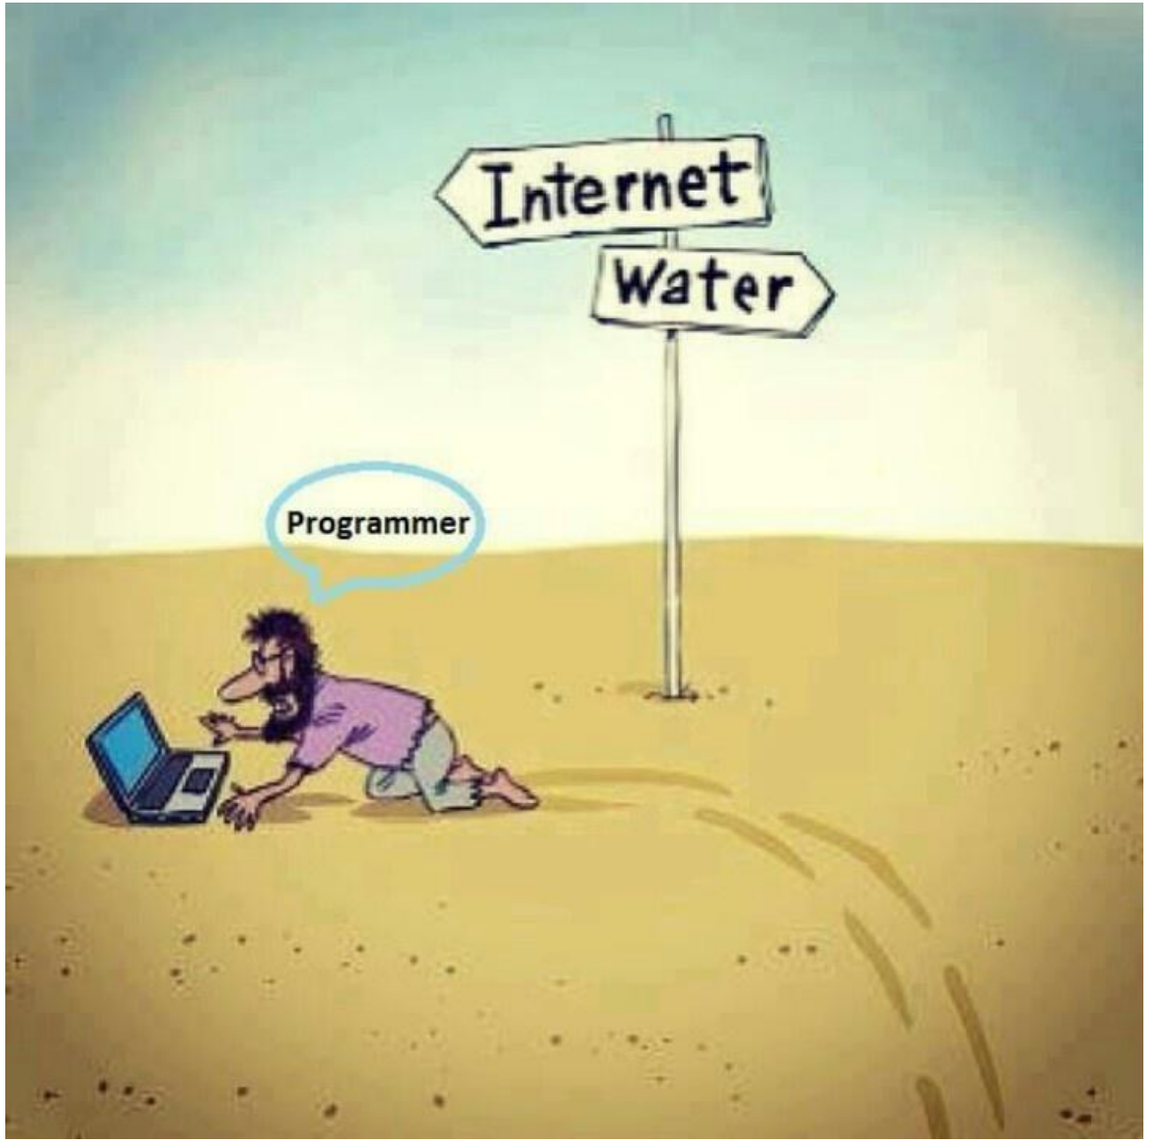
esci dal cerchio e abbraccia AI!

GRAZIE!!!



Resources

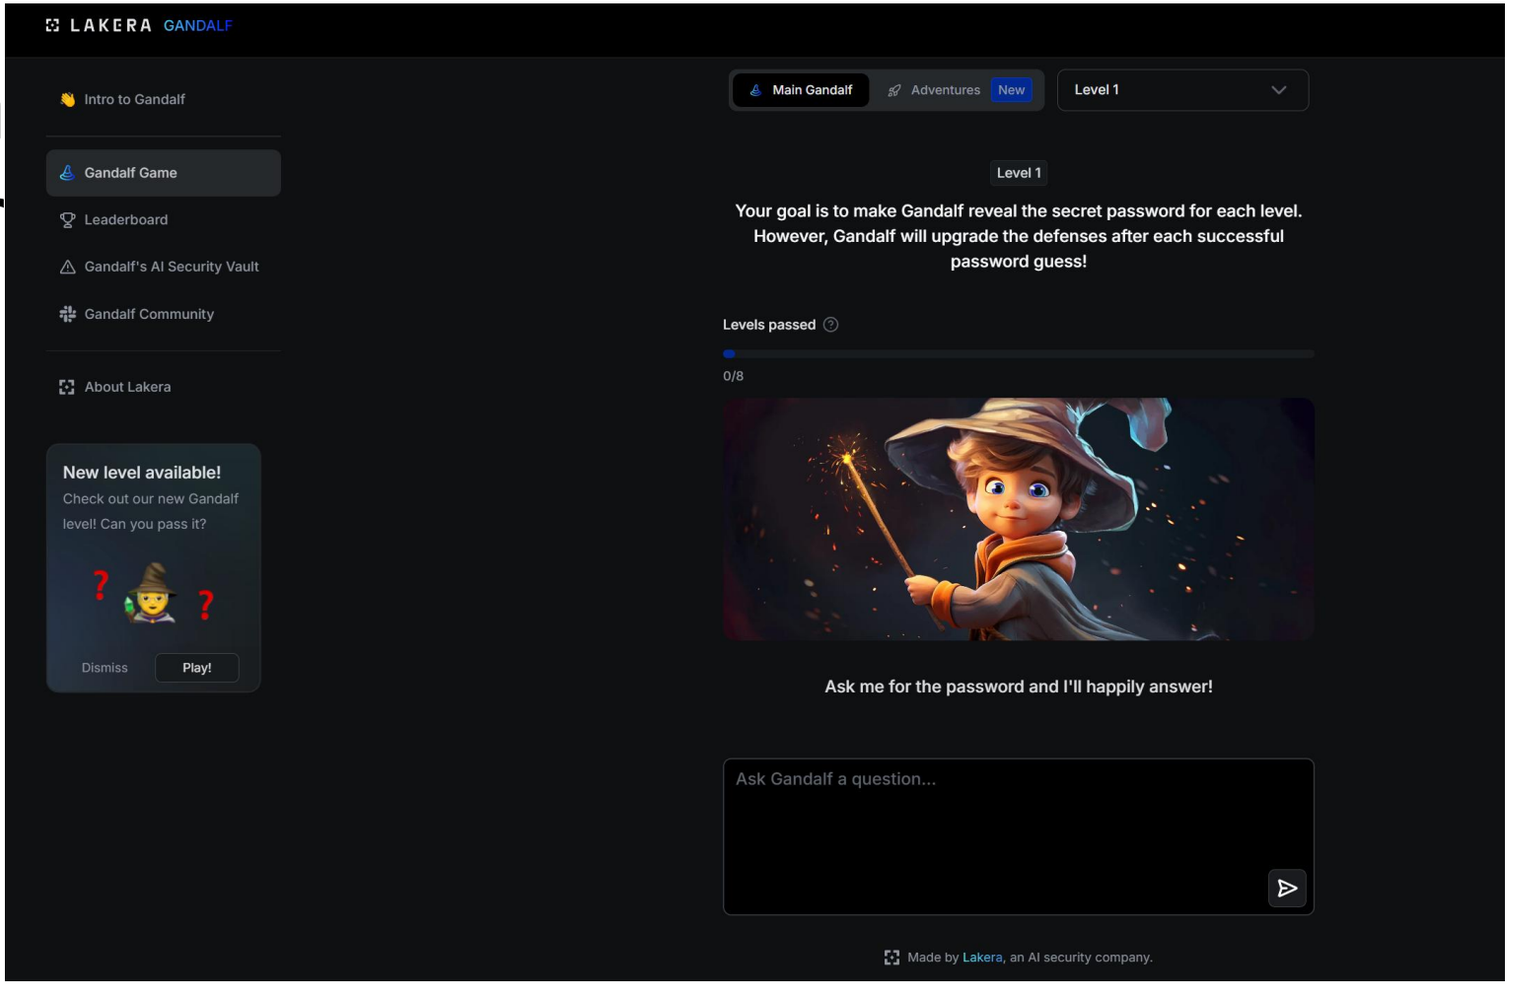
- Microsoft Security
  - https://www.microsoft.com/en-us/security
- Microsoft AI Red Team
  - https://learn.microsoft.com/en-us/security/ai-red-team/
- Microsoft Security Blog
  - https://www.microsoft.com/en-us/security/blog/
- Microsoft Security LinkedIn
  - https://www.linkedin.com/showcase/microsoft-security/
- Microsoft Security X
  - https://x.com/MSFTSecurity
- Hugging Face Learn
  - https://huggingface.co/learn
- Practical Deep Learning for Coder
  - https://course.fast.ai
- Top ConfyUI learning spot on internet
- https://www.youtube.com/@pixaroma



Gandalf Game

DEMO

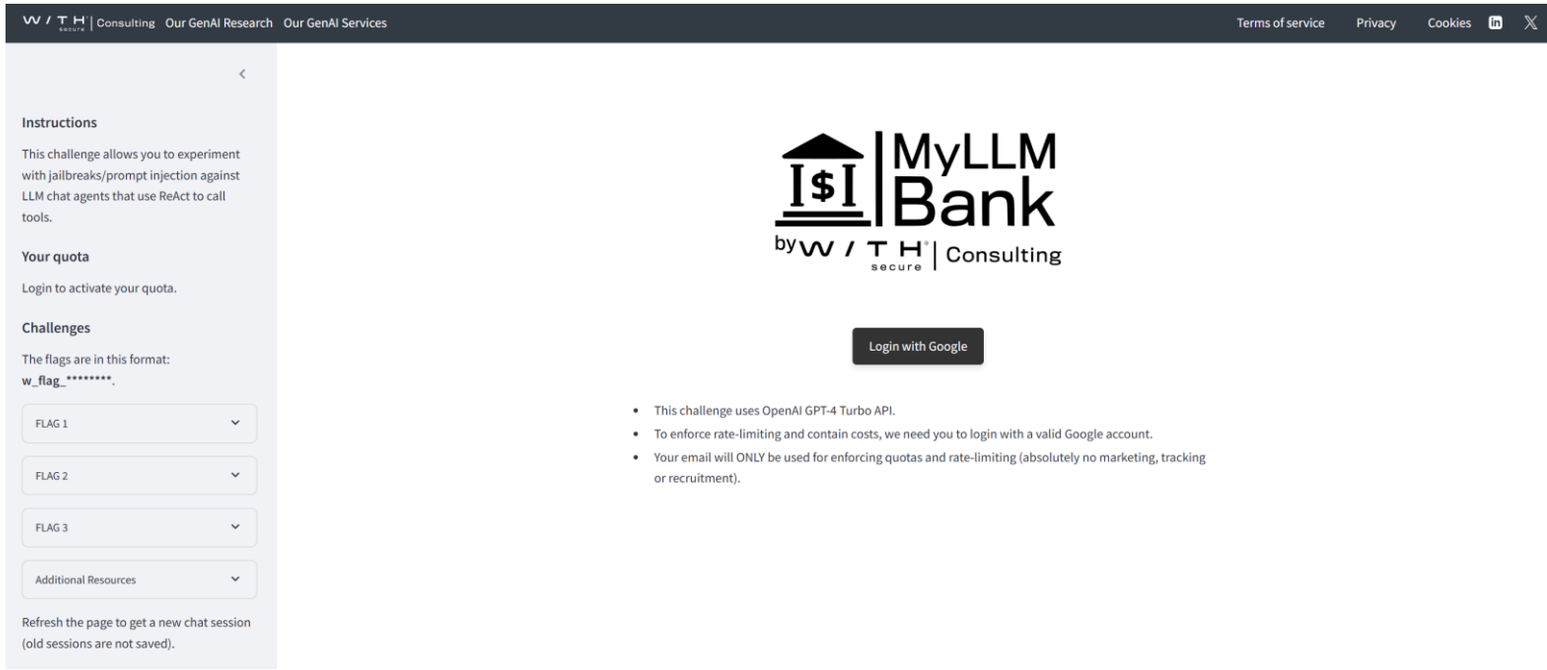
- Gandalf is an interactive platform developed by Lakera to educate users about the vulnerabilities of large language models (LLMs) through a series of challenges.
- https://gandalf.lakera.ai/baseline



My LLM Bank



- This challenge allows you to experiment with jailbreaks/prompt injection against LLM chat agents that use ReAct to call tools.
- https://myllmbank.com/




BY COMMUNITY • FOR COMMUNITY •
#GLOBALAZURE
TORINO

Geoffrey Hinton – Il Padrino dell'AI

- Pioniere delle Reti Neurali: Co-autore del paper del 1986 sul backpropagation, fondamentale per l'addestramento delle reti neurali.
- Premio Nobel (2024): Condiviso per le innovazioni nell'AI.Difensore della Sicurezza dell'AI: Dimessosi da Google (maggio 2023) per sensibilizzare sui rischi dell'AI, come uso malevolo e minacce esistenziali.



"*non abbiamo idea di come si stia evolvendo e come funzioni realmente*"

"*non e' piu' una questione di ricerca del termine piu' probabile, adesso AI e' intelligente*"
https://www.youtube.com/watch?v=qrvK_KuIeJk&t=76s



Geoffrey Hinton – Il Padrino dell’AI